

lab 2

2025-10-13

1 Examples

First, download “alr4” package.

```
options(repos = c(CRAN = "https://cloud.r-project.org/"))
install.packages("devtools")

## package 'devtools' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## D:\Temp\RtmpOsMlyF\downloaded_packages

library(devtools)

## Loading required package: usethis

devtools::install_github("cran/alr4")

## WARNING: Rtools is required to build R packages, but is not currently installed.
##
## Please download and install Rtools 4.5 from https://cran.r-project.org/bin/windows/Rtools/.

## Skipping install of 'alr4' from a github remote, the SHA1 (2250c241) has not changed since last install
## Use 'force = TRUE' to force installation
```

Example 1

“Heights” is a dataset consisting of height of mother “Mheight” and height of daughter “Dheight”. Assume the linear model: $dheight = a + b \times mheight + \epsilon, \epsilon \sim (0, \sigma^2)$.

```
install.packages("alr4")

## package 'alr4' successfully unpacked and MD5 sums checked
##
## The downloaded binary packages are in
## D:\Temp\RtmpOsMlyF\downloaded_packages
```

```

library(alr4)

## Loading required package: car

## Loading required package: carData

## Loading required package: effects

## lattice theme set by effectsTheme()
## See ?effectsTheme for details.

myfit=lm(dheight~mheight,data=Heights)
summary(myfit)

## 
## Call:
## lm(formula = dheight ~ mheight, data = Heights)
##
## Residuals:
##      Min    1Q Median    3Q   Max 
## -7.397 -1.529  0.036  1.492  9.053 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 29.91744   1.62247   18.44 <2e-16 ***
## mheight     0.54175   0.02596   20.87 <2e-16 ***  
## ---      
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 

## 
## Residual standard error: 2.266 on 1373 degrees of freedom
## Multiple R-squared:  0.2408, Adjusted R-squared:  0.2402 
## F-statistic: 435.5 on 1 and 1373 DF,  p-value: < 2.2e-16

coefficients(myfit)

## (Intercept)      mheight
## 29.917437     0.541747


```

Now we use scatter plot to check:

```

plot(Heights)
abline(myfit,col="red")


```



To see whether the model fits the data well, we check whether the residuals follow Gaussian distribution, and if there is a nonlinear relation between residuals and independent variables.


```
res=residuals(myfit)
qqnorm(res)#check if the residuals are normally distributed
```

Normal Q-Q Plot




```
hist(res)
```


Histogram of res





```
plot(Heights[, "mheight"], res)
```




```
plot(myfit)
```











```
## Example 2
```

Dataset “brains” provides brainweight(independent variable) and weights(response) of various animals.

```
#simple linear model
plot(brains[,2:1])
a=lm(BrainWt~BodyWt,data=brains)
abline(a)
```



```
#fit on log
plot(log(brains[,2:1]))
b=lm(BrainWt~BodyWt,data=log(brains))
abline(b,col=2)
```



We can see that fitting on log(variable) yields a better fit.

Now we predict the brain weight of an animal whose weight is 4kg.

```
x=as.numeric(coefficients(a)[1])
y=as.numeric(coefficients(a)[2])
x_log=as.numeric(coefficients(b)[1])
y_log=as.numeric(coefficients(b)[2])
bw=4
prediction=x+y*bw
print(prediction)
```

```
## [1] 94.87448
```

```
prediction_log=exp(x_log)*bw^y_log
print(prediction_log)
```

```
## [1] 23.97104
```

```
brains[brains[,2]<4.5&brains[,2]>3.5,]
```

	BrainWt	BodyWt
## Vervet	57.998	4.190
## Yellow-bellied marmot	17.000	4.050
## Rock hyrax2	21.000	3.600
## Raccoon	39.201	4.288
## Red fox	50.400	4.235


Apparently, the log prediction is closer to true data.

2 Exercise

2.1


We would like to find the relation between ranking and wealth(billion).

```
install.packages("readxl")  
  
## package 'readxl' successfully unpacked and MD5 sums checked  
  
## Warning: cannot remove prior installation of package 'readxl'  
  
## Warning in file.copy(savedcopy, lib, recursive = TRUE): problem copying  
## D:\360Downloads\Software\R-4.5.1\library\00LOCK\readxl\libs\x64\readxl.dll to  
## D:\360Downloads\Software\R-4.5.1\library\readxl\libs\x64\readxl.dll: Permission  
## denied  
  
## Warning: restored 'readxl'  
  
##  
## The downloaded binary packages are in  
## D:\Temp\Rtmp0sMlyF\downloaded_packages  
  
library(readxl)  
data=read_excel("forbes2019.xlsx")  
analysis <- data.frame(  
  ranking= data[[1]],  
  wealth= data[[4]]  
)  
plot(analysis)
```



From the plot, we can see that wealth and ranking cannot fit linear relation well. Therefore, we can first take log.

```
fit=lm(wealth~ranking,data=log(analysis))
plot(log(analysis))
abline(fit,col="red")
```



```

summary(fit)

##
## Call:
## lm(formula = wealth ~ ranking, data = log(analysis))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.34423 -0.03443 -0.02314  0.00228  0.21810
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 5.219425  0.036374 143.49 <2e-16 ***
## ranking     -0.568528  0.009708 -58.56 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.08949 on 98 degrees of freedom
## Multiple R-squared:  0.9722, Adjusted R-squared:  0.9719
## F-statistic: 3430 on 1 and 98 DF,  p-value: < 2.2e-16

```

```
coefficients(fit)
```

```

## (Intercept)      ranking
## 5.2194250 -0.5685282

```

Eventually, we can fit the data by $wealth = e^{5.22} ranking^{-0.57}$.

2.2

We would like to set up a standard for obesity.

(a) First, let's find the LS estimate for a, b and σ .

```
data <- read.table("height-weight.txt", header = TRUE)
df<-data.frame(
  sex=data[[1]],
  weight=data[[2]],
  height=data[[3]]
)

result <- lm(weight ~ height, data = log(df))

coef(result)

## (Intercept)      height
##   2.598779    2.929655

summary(result)

##
## Call:
## lm(formula = weight ~ height, data = log(df))
##
## Residuals:
##       Min     1Q   Median     3Q    Max 
## -0.27563 -0.07883 -0.00480  0.07677  0.45833 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 2.59878   0.08843   29.39   <2e-16 ***
## height      2.92966   0.16521   17.73   <2e-16 ***
## ---      
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1216 on 197 degrees of freedom
## Multiple R-squared:  0.6148, Adjusted R-squared:  0.6129 
## F-statistic: 314.4 on 1 and 197 DF,  p-value: < 2.2e-16
```

$\hat{a} = 2.598779, \hat{b} = 2.929655, \hat{\sigma} = 0.1216$.

```
a_hat=as.numeric(coefficients(result)[1])
b_hat=as.numeric(coefficients(result)[2])
sigma_hat=summary(result)$sigma
epsilon<-function(a_hat,b_hat,sigma_hat,W,H){
  return((log(W)-a_hat-b_hat*log(H))/sigma_hat)
}
is_obesity<-function(coefficient){
```

```

    return ((coefficient)>1.645)
}
#check my weight
W=73
H=1.78
my_coefficient=epsilon(a_hat=a_hat,b_hat=b_hat,sigma_hat=sigma_hat,W=W,H=H)
print(my_coefficient)

## [1] 0.01974951

print(is_obesity(my_coefficient))

## [1] FALSE

#check the proportion whose coefficient is larger than mine
proportion=pnorm(my_coefficient, lower.tail = FALSE)
print(proportion)

## [1] 0.4921216

```

My weight coefficient is about 0.02, and I have not exceeded the standard. There are around 50% people whose coefficients exceed me.

(b) If we directly regress on the data without taking log:

```

result2=lm(weight~height,data=df)
coef(result2)

## (Intercept)      height
## -130.7470     114.9222

summary(result2)

##
## Call:
## lm(formula = weight ~ height, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -19.650  -5.419  -0.576   4.857  42.887
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -130.747    11.563  -11.31  <2e-16 ***
## height       114.922     6.769   16.98  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.523 on 197 degrees of freedom
## Multiple R-squared:  0.594, Adjusted R-squared:  0.592 
## F-statistic: 288.3 on 1 and 197 DF,  p-value: < 2.2e-16

```

$$\hat{a} = -130.7470, \hat{b} = 114.9222, \hat{\sigma} = 8.523$$

```
a_hat2=as.numeric(coef(result2)[1])
b_hat2=as.numeric(coef(result2)[2])
sigma_hat2=summary(result2)$sigma
eta<-function(a,b,sigma,H,W){
  return((W-a-b*H)/sigma)
}
is_obesity2<-function(coefficient){
  return (coefficient>1.645)
}
#recheck my weight
W=73
H=1.78
my_coefficient2=eta(a=a_hat2,b=b_hat2,sigma=sigma_hat2,H=H,W=W)
print(my_coefficient2)
```

```
## [1] -0.09556973
```

```
print(is_obesity2(my_coefficient2))
```


```
## [1] FALSE
```

```
proportion=pnorm(my_coefficient2, lower.tail = FALSE)
print(proportion)
```


```
## [1] 0.5380688
```

The results from (a) and (b) are close. To see which one is more reasonable, we can check the plot of weight and height to see whether they follow a linear relation:

```
plot(df$height,df$weight)
```




```
plot(log(df$height),log(df$weight))
```




Still, the scatter plots look similar. Let's inspect the residuals of each regression:

```
#log regression  
plot(result,which=1)
```



```
#simple regression
plot(result2,which=1)
```



Unfortunately, it is still hard to distinguish which one is significantly better. However, since the R^2 of log linear regression is larger than simple linear regression, we should conclude that (a)log linear regression is slightly better.

Theoretically, using log linear regression is more reasonable, since weight and height ought to follow power law. Nevertheless, the reason why its performance is not significantly better than linear regression is probably that the range of height in the data is narrow.

(c) Now we add the factor of sex into the regression:

```
result3=lm(log(weight)~log(height)+sex,data=df)
coef(result3)
```

```
## (Intercept) log(height)          sex
##   3.0087053   2.0571556   0.1240784
```

We can see that $\hat{b} \approx 2$.

2.3

We would like to check Benford law by using Monte Carlo method.

```
set.seed(666)
#generate n samples
n=1000000
```

```

z=rnorm(n)
x=exp(z)
#find the first nonzero digit
transform<-function(x){
  return (x/10^floor(log10(x)))
}#transform the number such that the integer part is one digit
first_nonzero_digit <- function(x) {
  if (x == 0) return(0)

  transformed <- transform(x)

  first_digit <- floor(transformed)

  return(first_digit)
}
#result
first_digits <- sapply(x, first_nonzero_digit)
table(first_digits)/n

## first_digits
##      1       2       3       4       5       6       7       8
## 0.307945 0.170555 0.118611 0.093803 0.079311 0.068227 0.059750 0.053604
##      9
## 0.048194

i<-c(1,2,3,4,5,6,7,8,9)
Benford<-log10(1+1/i)
Benford

## [1] 0.30103000 0.17609126 0.12493874 0.09691001 0.07918125 0.06694679 0.05799195
## [8] 0.05115252 0.04575749

```

We can see that with n=1000000 samples, the simulation is very close to Benford law.